

Apache Hive – Best Practices

Mark Grover

Software Engineer, Cloudera Inc.

www.github.com/markgrover/bdtdc-hive

@mark_grover



Logistics

- Hadoop and Hive installed
 - Method 1: Demo VM at <https://bitly.com/PDfLKp>
 - Method 2: Install Hadoop + Hive (≥ 0.7)
- Dataset downloaded
- All instructions at github.com/markgrover/bdtdc-hive

Agenda

- Interactive
- Frequent breaks and conversations!
- Learn a little about Hive

My request to you

- Follow along
- Ask questions, any questions

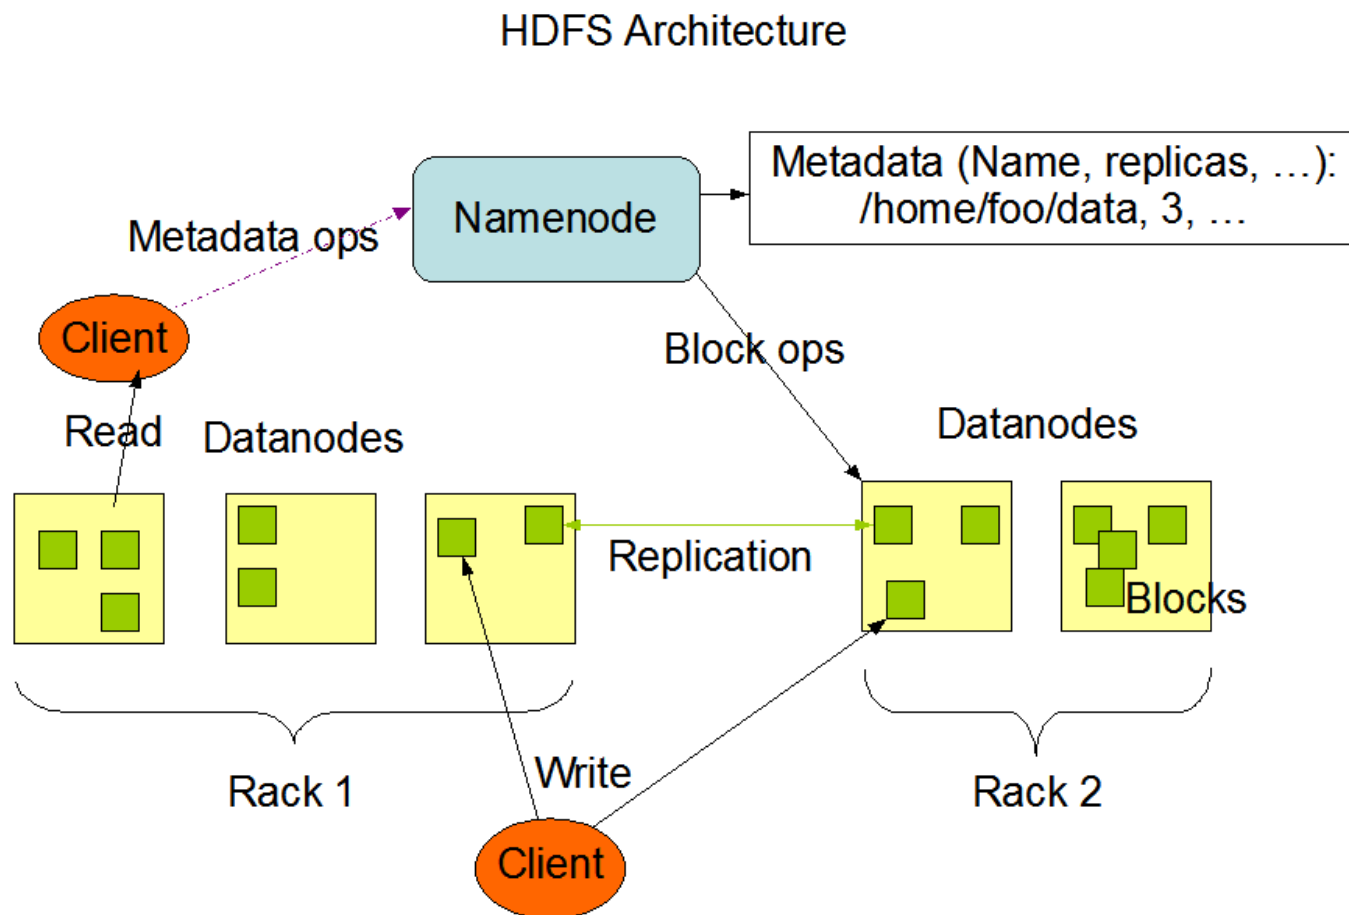
Hive

- What is Hive?
- SQL-like queries for processing Hadoop data

Hadoop

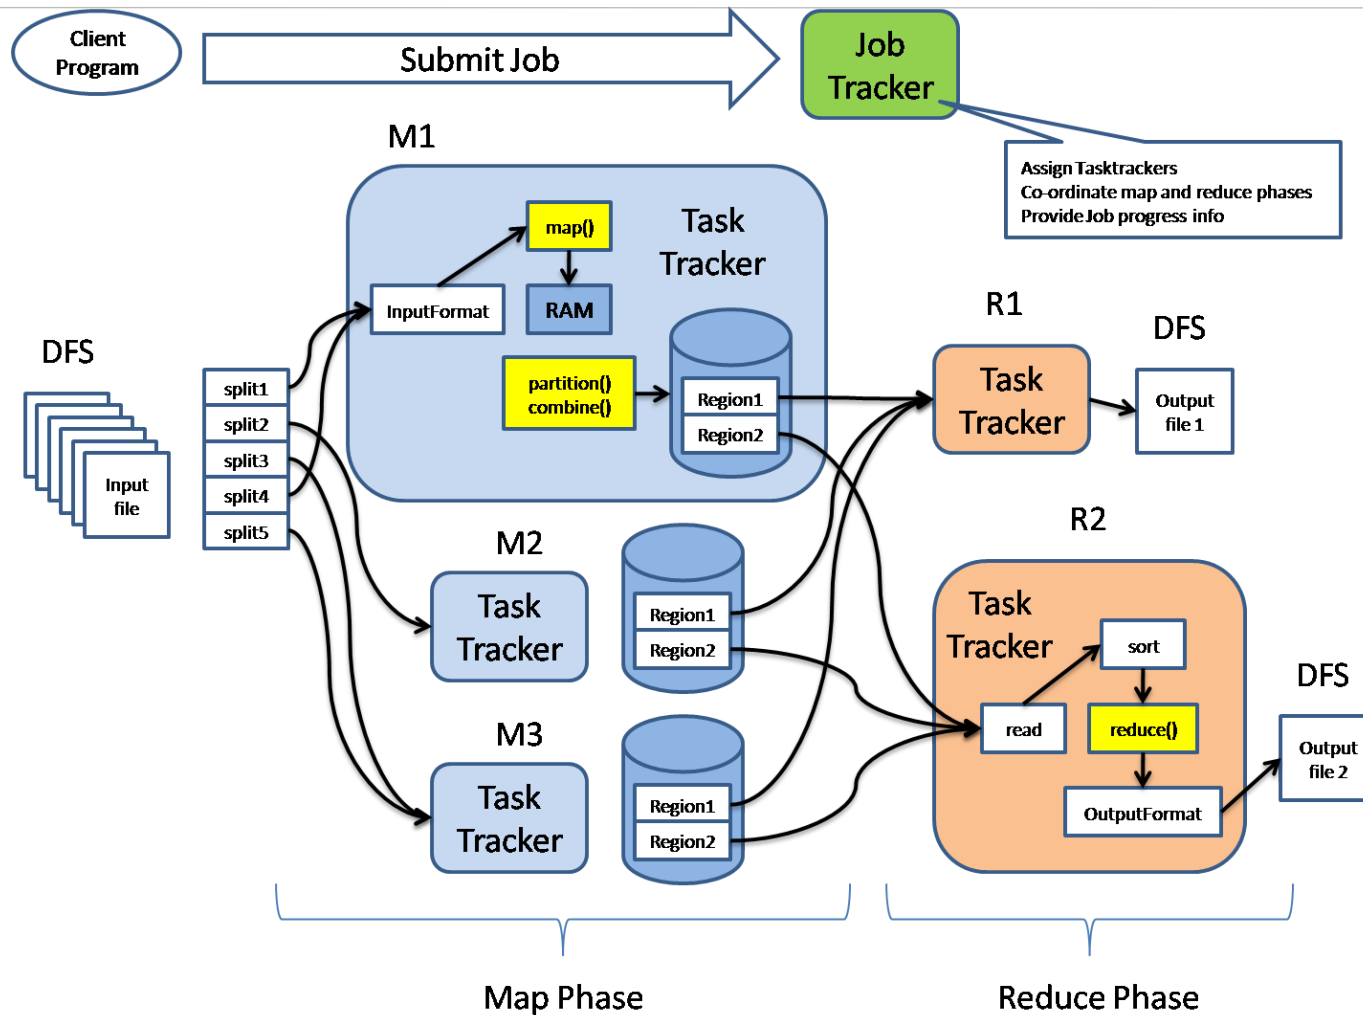
- What is Hadoop?
- Run some example MR jobs

HDFS Architecture



http://hadoop.apache.org/docs/stable/hdfs_design.html

How Hadoop (MR1) works



<http://architects.dzone.com/articles/how-hadoop-mapreduce-works>

The problem with MapReduce

- Catered towards developers
- No higher level abstraction
- Grouping, sorting, joining is very common but MR is tedious to write

Hive

- Why use Hive?
- Scalable
- Select, where clause, group by, order by, joins
- Pluggable User Defined Functions - UDFs (e.g. from_unixtime)
- Pluggable User Defined Aggregate Functions - UDAFs (e.g. count, avg)
- Pluggable User Defined Table Generating Functions - UDTFs (e.g. explode)

Hive...2

- Pluggable custom Input/Output format
- Pluggable Serialization Deserialization libraries (SerDes)
- Pluggable custom map/reduce scripts
- There is no HiveFormat!

What Hive does NOT currently support

- OLTP workloads - low latency
- Correlated subqueries

HiveFormat

- There is none!

HiveFormat

- There is none!
- One more time – there is none!
- Text files, SequenceFiles
- NoSQL (HBase, MongoDB)
- Custom input formats
- Virtual Input formats
 - Google docs!
 - SearchBotInputFormat

Let's get started

- Load the data on HDFS
- Create a Hive table
- Run some queries
- Connecting to Hive
 - Shell
 - JDBC
 - ODBC
 - Thrift client

Inside Hive

- Configuring Hive
- Hive metastore
- Partitioning
- Dynamic partitioning
- Bucketing
- Joins
- UDFs/UDAFs
- Custom map/reduce scripts
- Hive HBase Integration

Scalar Data types in Hive

- TINYINT
- SMALLINT
- INT
- BIGINT
- FLOAT
- DOUBLE
- DECIMAL (Note: Only available starting with Hive 0.11.0)
- TIMESTAMP (Note: Only available starting with Hive 0.8.0)
- DATE (Note: Only available starting with Hive 0.12.0)
- BOOLEAN
- STRING
- BINARY (Note: Only available starting with Hive 0.8.0)

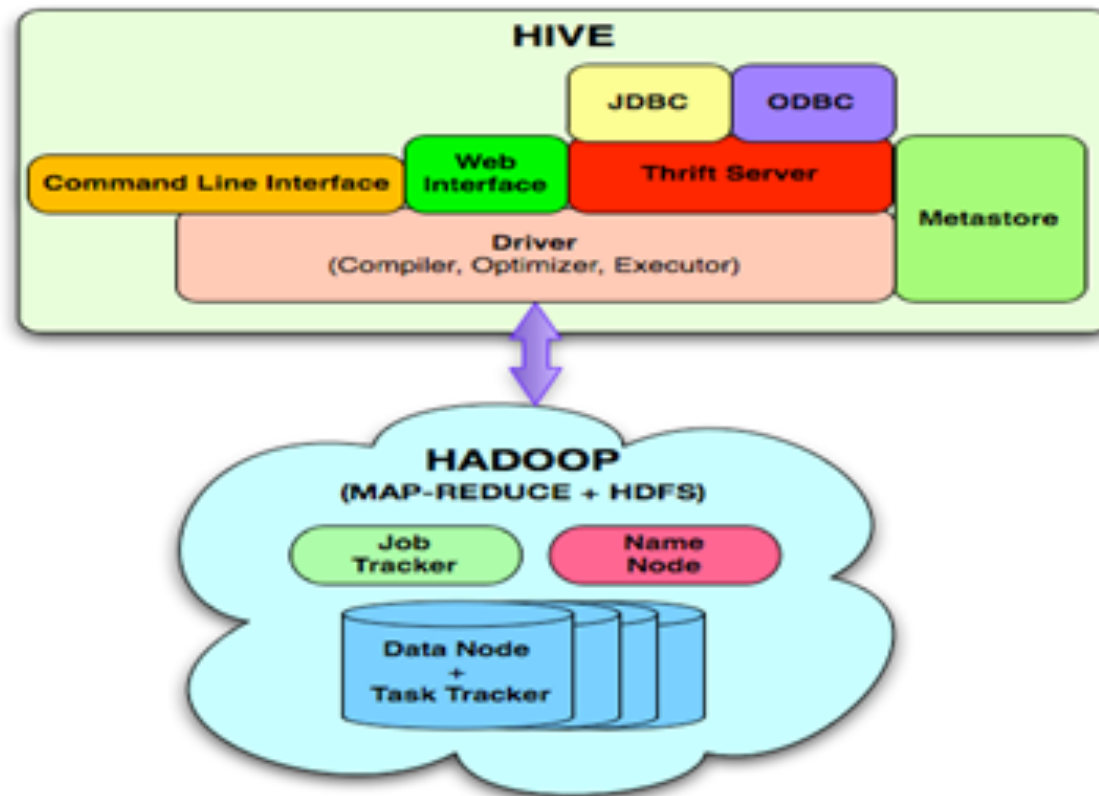
Complex data types in Hive

- Arrays
- Structs
- Maps
- Union

Configuring Hive

- hive-default.xml
- hive-site.xml (/etc/hive/conf/hive-site.xml)

Hive Architecture

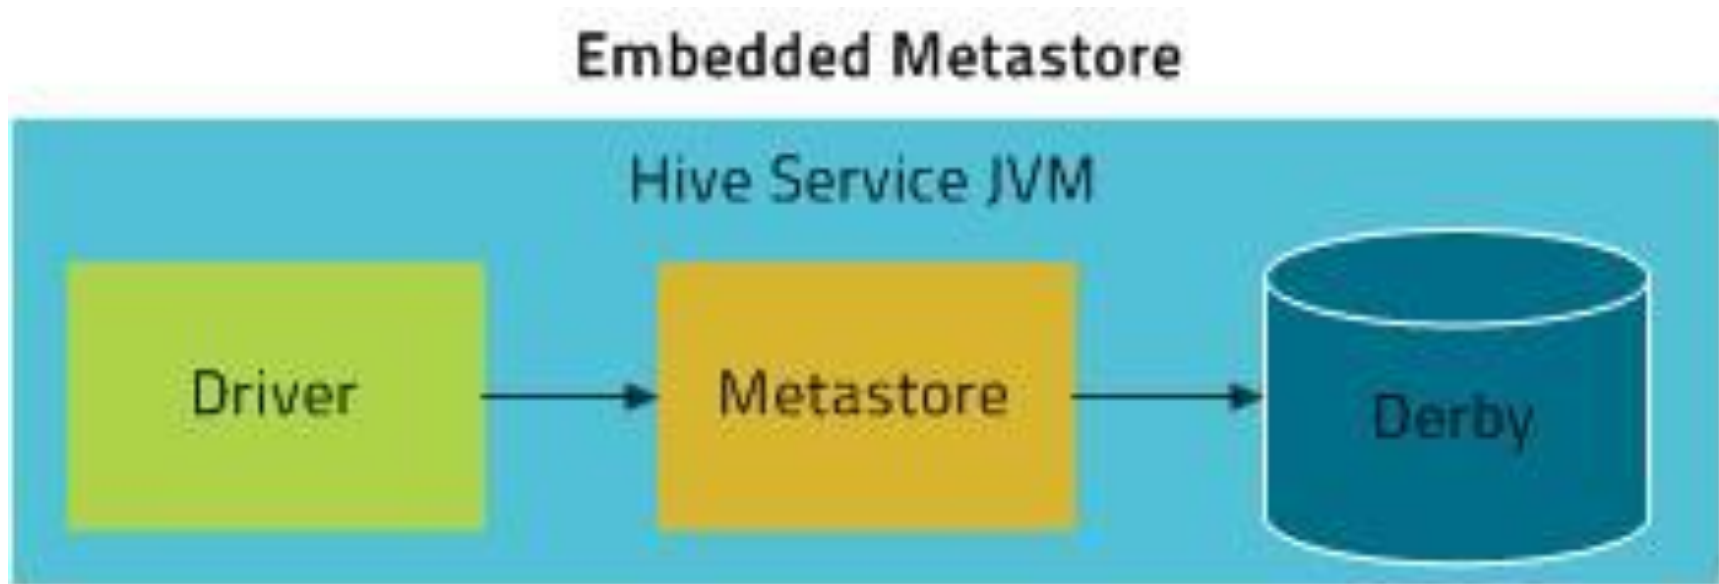


<http://okwang0921.blogspot.com/2012/11/hive-petabyte-scale-data-warehouse.html>

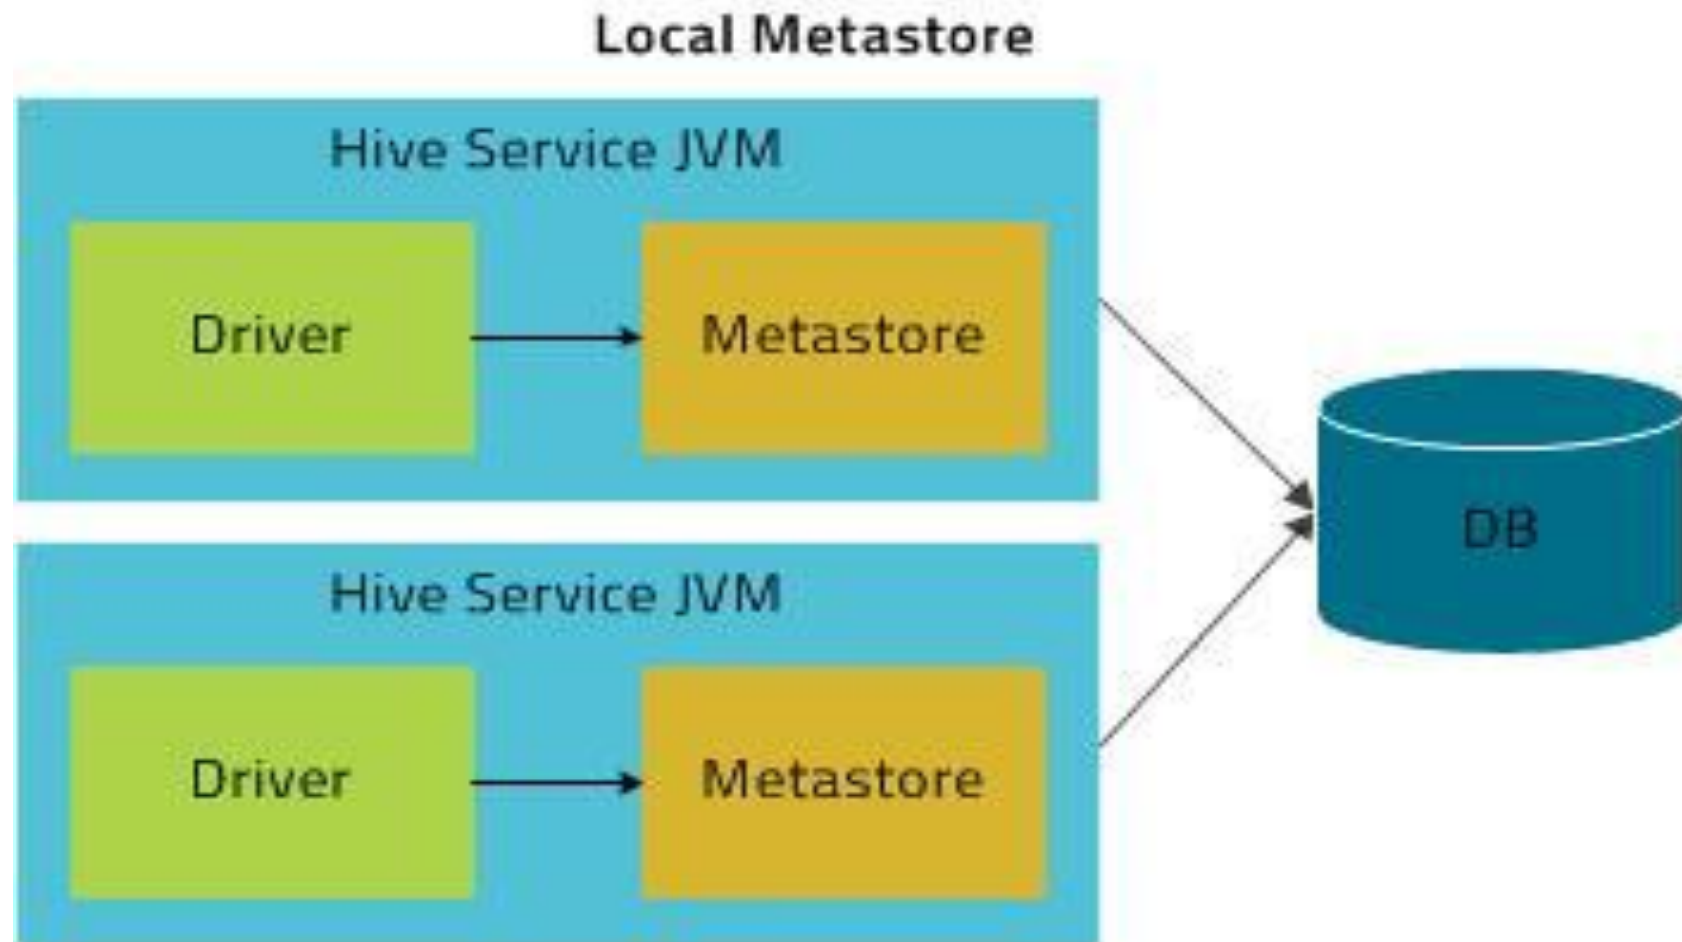
Hive metastore

- What is metastore?
- Vocabulary
 - Metastore
 - Metastore server (or datastore)
- Metastore service runs in different modes
 - Embedded
 - Local
 - Remote
- Default is embedded derby
 - Not recommended for anything non-trivial

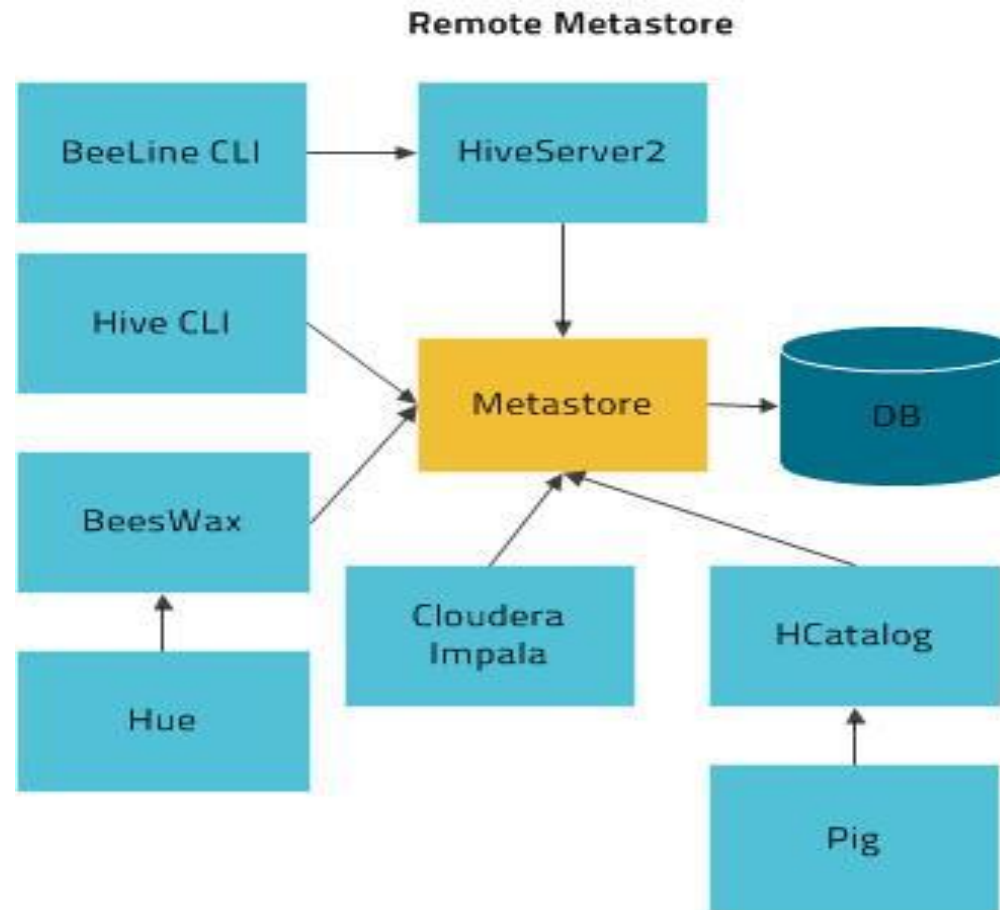
Embedded Metastore Server



Local Metastore Server



Remote Metastore Server



Configuring Hive metastore

- Default - embedded derby

<property>

<name>javax.jdo.option.ConnectionURL</name>

<value>jdbc:derby;;databaseName=metastore_db;create=true</value>

<description>JDBC connect string for a JDBC metastore</description>

</property>

Partitions in Hive

- Sub-directories under table directories
- Partition pruning

Partitions in Hive

- Create a partitioned table
- Run a query on the partitioned table and compare the time with non-partitioned table

Partitioning best practices

- Why/When to partition?
 - WHERE, GROUP BY, JOINS benefit
- Which column to partition by?
- How many partitions?
 - Don't overload the metastore (<10,000 partitions)
 - Don't run into small files problem

Dynamic partitioning

- Populate a partitioned table using Dynamic Partitioning

Joins in Hive

- Regular (aka common) joins
 - Reduce side
- Map joins
- Bucketed Map joins
- Sorted Bucketed Merge (SMB) joins

Joins: best practices

- When to use which join?
- If your data is sorted and bucketed
 - SMB join
- If your data is bucketed
 - Bucketed Map Join
- If all but one tables in the join can fit in memory of each node
 - Map join
- Regular join, otherwise

Bucketing in Hive

- Hash based bucketing within a partition
- For faster joins
- For sampling
- Sorted buckets

Sampling in Hive

- You only want a sample of the data
- Sampled on a particular column

Bucketing in Hive

- Create a bucketed table
- Perform a sampling query on this table

Bucketing best practices

- When to use bucketing?
- Which columns to bucket on?
- How many buckets to use?

UDFs/UDAFs

- What is a UDF?
- Custom Functionality
 - written in Java

Custom MR scripts

- Plug in custom MapReduce scripts in Hive queries

Custom plugin best practices

- When to use UDF and when to use custom MR scripts?
- Comparison between UDFs/UDAFs and custom MR scripts

Hive-HBase Integration

- Read from and write to HBase tables using Hive queries

Contact info

@mark_grover

github.com/markgrover

linkedin.com/in/grovermark

mgrover@cloudera.com

Don't Forget:

Ryan Blue's tutorial on Hadoop App development this afternoon!